

Introduction to Model Checking and NuSMV (2)

Cyrille Artho

KTH Royal Institute of Technology, Stockholm, Sweden
School of Electrical Engineering and Computer Science
Theoretical Computer Science

`artho@kth.se`

Outline

1. Liveness and safety.
2. How to use the NuSMV model checker, part 2.

Safety vs. Liveness

Safety: Something bad will never happen.

Ensures absence of defects and hazards.

Liveness: Something good eventually happens.

Ensures progress.

Counterexamples: Safety

Safety: Something bad will never happen.

Counterexample only needs to show one violation.




Counterexample is finite.

Counterexamples: Liveness

Liveness: Something good eventually happens.

Counterexample needs to show that something never happens.




Counterexample is infinite.

Safety vs. Liveness

Safety: Something bad will never happen.

Ensures absence of defects and hazards.


Counterexample is a **finite** trace.



Liveness: Something good eventually happens.

Ensures progress.

Counterexample is an **infinite** trace.



Which temporal logic operators are suitable for which type of property?

next (var) VS. X

next (var) : refer to **value of a given variable in the next state**.

You cannot use "X" in this case, because X refers to the value of a *predicate* in the next state.


Therefore, **X (n_items) = n_items - 1** will not work, because it will compare n_items with itself in the next state, i.e., n_items - 1 will also refer to the value of n_items in the next state.

To compare a value in the current and next state, use **next** in an LTL formula:

next (n_items) = n_items - 1

This will use the future value of n_items on the left side and the current value on the right side.

Applying NuSMV: Ferryman puzzle



- Ferryman wants to cross river with cabbage (c), goat (g), wolf (w).
- Goat will eat cabbage when left alone; wolf will eat goat.
- Ferry carries only one „passenger”.

Can the ferryman bring all things to the other side, safely?

Wolf, goat, ferryman, river icons made by Freepik; cabbage icon made by Nikita Golubev; icons from www.flaticon.com

NuSMV model of the ferryman puzzle state

```
-- Ferryman by Bow-Yaw Wang
MODULE main
VAR
    ferryman : boolean;
    goat : boolean;
    cabbage : boolean;
    wolf : boolean;
    carry : { g, c, w, n };

ASSIGN
    init (ferryman) := FALSE;
    init (goat) := FALSE;
    init (cabbage) := FALSE;
    init (wolf) := FALSE;
    init (carry) := n;
```

- Boolean variables **ferryman**, **goat**, **cabbage**, **wolf** denote the location of the ferryman, goat, cabbage, wolf.
- Initially, all are on the same side (**FALSE**).
- The variable **carry** denotes the good carried by the ferryman: **g** (goat), **c** (cabbage), **w** (wolf), or **n** (none).

Modeling the ferryman and his passengers

```
next (ferryman) := { FALSE, TRUE };
next (goat) := case
    ferryman = goat & next (carry) = g: next (ferryman);
    TRUE: goat;
esac;
next (cabbage) := case
    ferryman = cabbage & next (carry) = c: next (ferryman);
    TRUE: cabbage;
esac;
next (wolf) := case
    ferryman = wolf & next (carry) = w: next (ferryman);
    TRUE: wolf;
esac;
```

- The ferryman is non-deterministic (we don't know the right strategy).
- The passengers follow the ferryman iff he carries them to other side.

Modeling the possibility to carry a passenger

TRANS

```
(next (carry) = n) |  
(ferryman = goat & next (carry) = g) |  
(ferryman = cabbage & next (carry) = c) |  
(ferryman = wolf & next (carry) = w);
```

- The ferryman can carry nothing, or...
- starts from the same place as the item he will carry in the next turn.
- **TRANS** is another way to model transition relation.

How to describe the puzzle and its solution?

- Remember: ferryman needs to watch if goat is together with cabbage or wolf.
- Therefore: if goat is on the same side as cabbage or wolf, ferryman must be on that side, too.
- Once all four are on the other side, puzzle is solved.
- Use **until** operator:

*rules of puzzle are followed **U** solution achieved.*

Encoding the puzzle in LTL

$((\text{goat} = \text{cabbage} \mid \text{goat} = \text{wolf}) \rightarrow \text{goat} = \text{ferryman})$ *[rule]*



$\mathbf{U} (\text{cabbage} \ \& \ \text{goat} \ \& \ \text{wolf} \ \& \ \text{ferryman})$ *[solution]*

We want to see the solution!

We negate the whole property, stating
„I can't follow the puzzle rules **until** the solution is achieved”.

$\mathbf{!} ((\text{goat} = \text{cabbage} \mid \text{goat} = \text{wolf}) \rightarrow \text{goat} = \text{ferryman})$
 $\mathbf{U} (\text{cabbage} \ \& \ \text{goat} \ \& \ \text{wolf} \ \& \ \text{ferryman}))$

Counterexample = solution



Another bridge crossing puzzle: „Bridge and torch problem”

- A, B, C, and D want to cross a bridge at night.
- They have only one weak torch that gives light for up to two people.
- Torch must be carried across the bridge (cannot be thrown across).
- The time taken for each crossing depends on the slowest person:

A	1
B	2
C	5
D	10

- Can all four cross within 17 time units?

An example crossing

Left side				Right side	Time
A	B	C	D		0
				A B	2
A		C	D	B	3
			D	A B C	8
A			D	B C	9
				A B C D	19

Can we do better?

NuSMV model: Variables

- Location of A, B, C, D are booleans (or an array of booleans).
- Another array of booleans denote of A, B, C, D are *traveling*.
- Torch location is also a boolean.
- Time is a number between 0 and 100.

Some transitions

- Torch can change location only if someone travels.
Also possible to model that torch always changes location until solution achieved, since we are interested in efficient solutions, and time does not increase if nobody moves.
- Choice of who travels is not specified.
- Location of A, B, C, D is updated iff
 1. they want to travel,
 2. the torch is at their place.

Timekeeping

- Time advances according to the slowest person who travels.
If you model torch moves as optional (see above), then ensure that „empty moves” do not count towards time, by incrementing time only if `location(x) != next(location(x))`;
- You need to have a non-overflow rule at top, e. g.,

```
next(time) := case  
  time > 90: 90;
```

The final formula and game rule

- At most two people travel: use **count**:
count (. . . , . . . , . . .)
returns the number of true predicates in the list.
- A, B, C, D have to arrive at the other side within N time units.
- Use !, “follow game rules **U** goal” as template;
solution is counterexample.
- You can add other consistency checks, such as
G (**location**[0] & **location**[1] & **location**[2] & **location**[3]
→ **torch**)
- The time limit is another conjunct in the goal condition.
- What happens if you lower the time limit further?

Lab exercise 2 (in assignment 1)

1. Develop the full game model and the formula to find the solution.
2. Study the resulting trace(s) and show the optimal solution.

Semaphores



Photo by Dave F

- Inspired by railway signal.
- Binary semaphore:
Resource can be available or in use.
- Value of semaphore guards access to exclusive shared resource („critical section” in concurrent code).
- Anyone who wants access to resource in use, has to wait.

Semaphore model

```
MODULE user(semaphore)
VAR
  state : { idle, entering, critical, exiting };
ASSIGN
  init(state) := idle;
  next(state) := case
    state = idle : { idle, entering };
    state = entering & !semaphore : critical;
    state = critical : { critical, exiting };
    state = exiting : idle;
    TRUE : state;
  esac;
  next(semaphore) := case
    state = entering: TRUE;
    state = exiting : FALSE;
    TRUE : semaphore;
  esac;
```

- Process („user”) may only acquire semaphore when not in use.

Semaphore model – 2

```
MODULE main
VAR
  semaphore : boolean;
  proc1 : process user(semaphore);
  proc2 : process user(semaphore);
ASSIGN
  init(semaphore) := FALSE;
SPEC -- safety
  AG !(proc1.state = critical & proc2.state = critical);
SPEC -- liveness
  AG (proc1.state = entering -> AF proc1.state = critical);
```

- System is *asynchronous*: only one process runs at a time!
- Therefore, semaphore is only updated by active process; this ensures that no two processes can obtain it at once.

Liveness property does not hold!

`AG (proc1.state = entering -> AF proc1.state = critical)`
is **false**

Trace Type: Counterexample

```
-> State: 1.1 <-  
    semaphore = FALSE  
    proc1.state = idle  
    proc2.state = idle  
-> Input: 1.2 <-  
    _process_selector_ = proc1  
    running = FALSE  
    proc2.running = FALSE  
    proc1.running = TRUE  
-- Loop starts here  
-> State: 1.2 <-  
    proc1.state = entering  
-> Input: 1.3 <-  
    _process_selector_ = proc2  
    proc2.running = TRUE  
    proc1.running = FALSE  
-> State: 1.3 <-
```

- Trace shows interleaving between processes.
- Process 1 is selected, tries to acquire semaphore.
- After that, process 2 is always selected.
- Process 1 never gets another turn.

This is not fair!

- A good „scheduler” would always eventually run process 1.
- Error traces from unfair runs are not always relevant.
- We can forbid such scenarios with *fairness constraints*:

FAIRNESS `running`;

- Only traces where `running` is true infinitely often are considered.

It is possible to list multiple fairness conditions!

Another error trace

```
-> State: 1.1 <-  
    semaphore = FALSE  
    proc1.state = idle  
    proc2.state = idle  
-> Input: 1.2 <-  
    _process_selector_ = proc1  
    running = FALSE  
    proc2.running = FALSE  
    proc1.running = TRUE  
-- Loop starts here  
-> State: 1.2 <-  
    proc1.state = entering  
-> Input: 1.3 <-  
    _process_selector_ = proc2  
    proc2.running = TRUE  
    proc1.running = FALSE  
-- Loop starts here  
-> State: 1.3 <-  
-> Input: 1.4 <-
```

```
-> State: 1.4 <-  
    proc2.state = entering  
-> Input: 1.5 <-  
-> State: 1.5 <-  
    semaphore = TRUE  
    proc2.state = critical  
-> Input: 1.6 <-  
    _process_selector_ = proc1  
    proc2.running = FALSE  
    proc1.running = TRUE  
-> State: 1.6 <-  
-> Input: 1.7 <-  
    _process_selector_ = proc2  
    proc2.running = TRUE  
    proc1.running = FALSE  
-> State: 1.7 <-  
    proc2.state = exiting  
-> Input: 1.8 <-  
-> State: 1.8 <-  
    semaphore = FALSE  
    proc2.state = idle
```

What goes wrong here?

Analyze the error trace with your lab partner(s).

- What happens in this error trace?
- What is the problem (the infinite loop)?
- Is the fairness condition fulfilled?
- What could be changed to avoid such behavior?

Processes in NuSMV

- Used to model systems where only one component is active at a time.
- No „global clock” (synchronous behavior): *asynchronous systems*.
- Problem:

WARNING * Processes are still supported, but deprecated. *****
WARNING * In the future processes may be no longer supported. *****

Write your own model code to “schedule” modules.

How to model processes with variables

- NuSMV selects the active process in between regular model steps.
- We need a variable that holds the ID of the active process:

VAR running: 0..N;

- We cannot interleave the choice of „running” with other modules!

Use next (running) = PID as scheduling predicate.

Lab exercise 3 (assignment 1): Change in semaphore model

1. User module

```
MODULE user(semaphore, active)  
  -- state transitions only fire when active
```

2. Semaphore module

- Semaphore is now executed globally, once per turn.
- Condition needs to be repeated for each user instance, but...
state only changes if condition is true for **active process!**

Adapt semaphore model and update the fairness conditions.

Guidance for exercise 3

1. Transition rule of module „user” easy to adapt:
State changes only if „active” is true.
2. Process scheduling:
`running: 0..1;`
`proc1 : user(semaphore, next(running) = 0);`
`proc2 : user(semaphore, next(running) = 1);`
3. **Semaphore transition cannot be specified inside „user”!**
Reason: Semaphore is executed globally, only once per step;
so it has to be specified globally.
 - (a) Copy/paste transition function; adapt for each user process.
 - (b) Transition has to take into account „active” flag of each process.
4. Fairness condition: Each process has to run infinitely often.
More fairness conditions may be needed to avoid livelock, fulfill properties.

Assignment 1: Three exercises


1. Vending machine: fix, add counter.
 2. Bridge crossing: implement, find optimum.
 3. Semaphore: find solution without NuSMV processes.
- All parts: Add brief comments (starting with `--`) for variables, modules, transitions, properties.
 - What does variable represent, why is starting state as is.
 - What does transition rule/property mean, why was it changed.
 - Part 2: Add explanation of optimum (error trace) as comment at end. Multi-line comments start with `/--`, end with `--/`.
 - Part 3: Add explanation of error trace when fairness is disabled.

Lab exercise

Lab: opportunity to ask questions and work on exercise.

1. Introductory exercise:
Submit on Canvas before Friday! (Imperfect solution ok)
2. Large (graded) exercise in labs 2 and 3 of each cycle.


Model Checking = state space search



- How to analyze reachable states against the property?

How to check a model against a property

- Properties can be translated into so-called Büchi automata.
- These automata can be non-deterministic and accept infinite words.
- Automata can be determinized (at exponential cost).



• $0 \in L(A)$?

• 01?

• 00?


• 10?

• 100?

• 101?

How to implement a model checker

- **Negate** property $\neg p$: **error state(s)**
(accepting states in Büchi automaton).
- Cross-product of model with that automaton encoding $\neg p$.
- All we need to do now is to check whether „bad” states are reachable.



SPIN

- Popular explicit-state open source model checker.
- Open source since 1991.
- Compiles transition system (in Promela) to C code.
- Uses optimizations such as state hashing to speed up search.
- Very fast if entire model fits in RAM.

What if model is too big to fit in memory?

Binary Decision Diagrams (BDDs)

- Published by Randal Bryant in 1986.
- Compressed representation of Boolean functions.
- Eliminates redundancy in sub-expressions.
- For equal variable order, reduction is **canonical** (same result regardless of order in which subexpressions are added).
- Boolean **operations** can be implemented on **compressed** data (ROBDDs).